

week4_code

February 11, 2025

1 Week 2 Ingesting and Exploring the Dataset

```
[1]: # install wordcloud
!pip install wordcloud
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: wordcloud in /home/jupyter-
geean/.local/lib/python3.12/site-packages (1.9.4)
Requirement already satisfied: numpy>=1.6.1 in
/opt/tljh/user/lib/python3.12/site-packages (from wordcloud) (1.26.4)
Requirement already satisfied: pillow in /opt/tljh/user/lib/python3.12/site-
packages (from wordcloud) (11.1.0)
Requirement already satisfied: matplotlib in /opt/tljh/user/lib/python3.12/site-
packages (from wordcloud) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud) (1.3.1)
Requirement already satisfied: cycler>=0.10 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud)
(0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud)
(4.55.3)
Requirement already satisfied: kiwisolver>=1.3.1 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud) (1.4.8)
Requirement already satisfied: packaging>=20.0 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud) (24.1)
Requirement already satisfied: pyparsing>=2.3.1 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/tljh/user/lib/python3.12/site-packages (from matplotlib->wordcloud)
(2.9.0.post0)
Requirement already satisfied: six>=1.5 in /opt/tljh/user/lib/python3.12/site-
packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.17.0)
```

```
[2]: # import packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

from wordcloud import WordCloud
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")

```

```

[3]: # change working directory
import os
os.getcwd() # Get current working directory
os.chdir('..') # Move up one directory level from notebooks
print(os.getcwd())
#os.chdir('../data') # change to the data folder

```

/home/jupyter-geeant/cookiecutter-data-science/{{ cookiecutter.repo_name }}

```

[4]: # load the data
df = pd.read_csv('data/Combined Data.csv', index_col=0)

```

```

[5]: # make a copy and get rid of the missing values
df1 = df.copy()
df1.dropna(inplace = True)
# see the top head of the data
df1.head()

```

```

[5]:

```

	statement	status
0	oh my gosh	Anxiety
1	trouble sleeping, confused mind, restless hear...	Anxiety
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety
3	I've shifted my focus to something else but I'...	Anxiety
4	I'm restless and restless, it's been a month n...	Anxiety

```

[6]: # number of missing values
missing_values = df.isnull().sum()

print(missing_values)

```

```

statement    362
status        0
dtype: int64

```

```

[7]: # get the rows and columns of all of the data
rows,columns = df.shape
print(f"Number of rows: {rows}")
print(f"Number of columns: {columns}")

```

```

Number of rows: 53043
Number of columns: 2

```

```
[8]: # calculate the number of missing values
rows_with_missing = df[df.isnull().any(axis=1)]
print(rows_with_missing)
```

	statement	status
293	NaN	Anxiety
572	NaN	Anxiety
595	NaN	Anxiety
1539	NaN	Normal
2448	NaN	Normal
...
52838	NaN	Anxiety
52870	NaN	Anxiety
52936	NaN	Anxiety
53010	NaN	Anxiety
53031	NaN	Anxiety

[362 rows x 2 columns]

The dataset contains 362 missing values in the ‘Statement’ column and no missing values for ‘Status’.

2 Missing Values -Week 3

Many of the rows have NaNs and represent anxiety and normal. Since there are 53,043 values and there are only 362 rows where there is missing values. We feel that it is best to drop these rows since they represent only 0.7% of the data and as you will see later we have an abundance of “normal” and “anxiety” labeled data.

The dataset includes 52,681 rows and 2 columns after removing missing values.

```
[9]: # get the rows and columns of the data that drops the missing values
rows,columns = df1.shape
print(f"Number of rows: {rows}")
print(f"Number of columns: {columns}")
```

Number of rows: 52681

Number of columns: 2

We want to add a column to explore the length of each statement. This can help us quantify the user’s input and support further analysis. This will give us an idea on how to preprocess the text and determine tokenization especially for transformer models. Many NLP models, especially those based on deep learning, have limitations on input length so determining the length is important.

```
[10]: # create a new column that gives the length of each statement
df1['statement_len'] = df1['statement'].apply(lambda x: len(x.split(' ')))
df1.head()
```

```
[10]:
```

	statement	status	statement_len
0	oh my gosh	Anxiety	3
1	trouble sleeping, confused mind, restless hear...	Anxiety	10
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14
3	I've shifted my focus to something else but I'...	Anxiety	11
4	I'm restless and restless, it's been a month n...	Anxiety	14

From the output, we can see that this dataset includes 2 variables: statement and status.

The statement variable is a text variable that contains different user inputs.

The status variable represents different emotional statuses, which contain different categories.

The next step is to explore dataset

```
[11]: # information about the dataset
'''The class type of the DataFrame.
The range of the index.
The number of columns and their names.
The count of non-null values in each column.
The data type of each column.
The memory usage of the DataFrame.'''

print(df1.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 52681 entries, 0 to 53042
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   statement       52681 non-null  object
1   status          52681 non-null  object
2   statement_len   52681 non-null  int64
dtypes: int64(1), object(2)
memory usage: 1.6+ MB
None
```

Statement and status column are object data types. The statement_len column is an integer/numeric datatype.

```
[12]: # descriptive statistics
'''count is the number of non-null entries.
unique is the number of unique values.
top is the most frequent value.
freq is the frequency of the most frequent value.'''
df1.describe(include='object').T
```

```
[12]:
```

	count	unique	top	freq
statement	52681	51073	what do you mean?	22
status	52681	7	Normal	16343

The 'Statement' column contains 51,073 unique values, indicating that most user inputs are unique. The most frequently appeared statement is "What do you mean?" and occurred 22 times in the dataset. The frequent occurrence of "What do you mean?" suggests significant communication gaps or misunderstandings, indicating areas where individuals feel confused or need more clarity, which is crucial in mental health discussions. This phrase often reflects a state of uncertainty or anxiety, signaling important emotional states. It could also indicate active engagement and a desire for better understanding and it could indicate that individuals need more support or reassurance, aiding in tailoring mental health resources effectively.

The 'Status' column contains 7 unique values and represents different emotion statuses. The most common status is "Normal", suggesting that over 30% of the statements in the dataset fall under this category.

```
[13]: # Get summary statistics for the 'statement_len' column
summary_statistics = df1['statement_len'].describe()
print(summary_statistics)
```

```
count      52681.000000
mean        113.035914
std         163.501877
min          1.000000
25%         15.000000
50%         62.000000
75%        148.000000
max        6300.000000
Name: statement_len, dtype: float64
```

```
[14]: # Calculate the mode of the 'statement_len' column
mode_value = df1['statement_len'].mode()[0]

print(f"The mode of the 'statement_len' column is: {mode_value}")
```

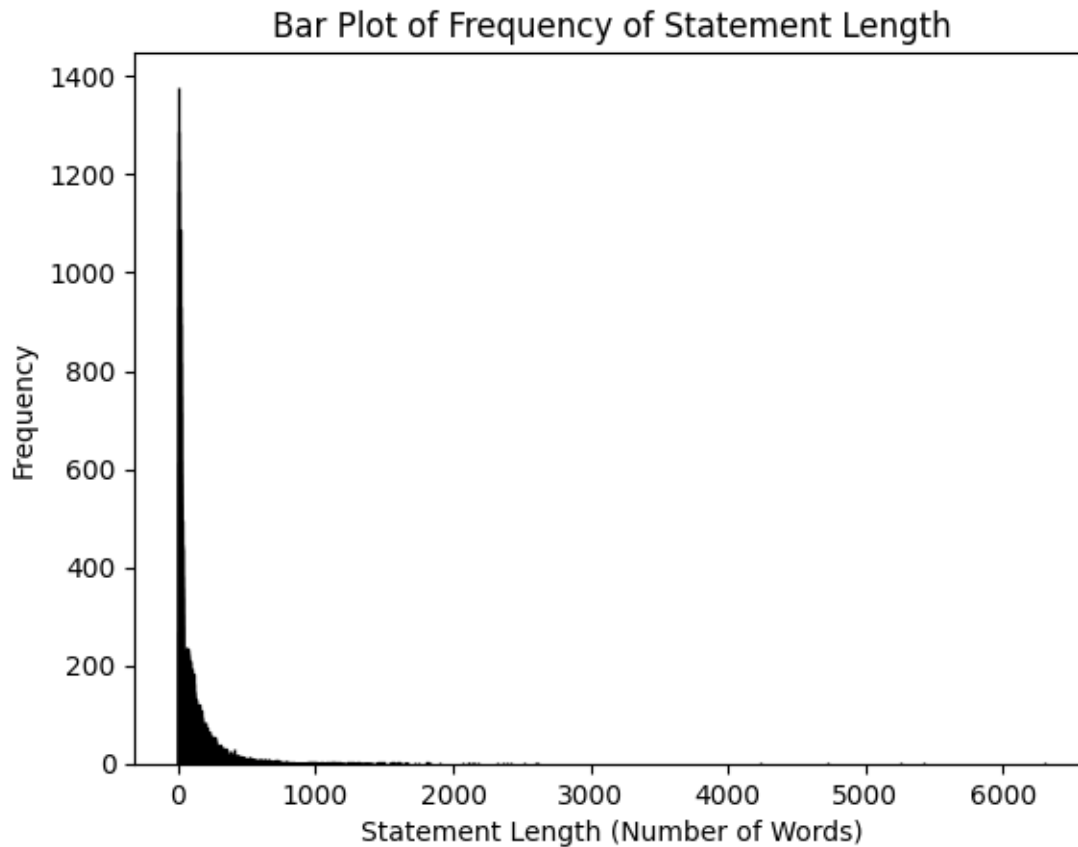
The mode of the 'statement_len' column is: 5

The summary statistics for the 'Statement_len' column show the distribution of statement lengths. The average statement contains 113 words with a standard deviation of 163.5 words. The shortest statement only has 1 word, while the longest contains 6300 words. The most frequent statement length is 5 words, indicating that short phrases are commonly used.

The following bar plot of the frequency of statement length visualizes the previous statement.

```
[15]: # Create a bar plot of the frequency of the 'statement_len' column
statement_len_counts = df1['statement_len'].value_counts()

plt.bar(statement_len_counts.index, statement_len_counts.values,
        edgecolor='black')
plt.xlabel('Statement Length (Number of Words)')
plt.ylabel('Frequency')
plt.title('Bar Plot of Frequency of Statement Length')
plt.show()
```



The histogram shows that it is a right skewed distribution, which most of the statement length under 1000 words. This means that when we focus on the output length, we should set it to be under 1000.

```
[16]: # Histogram of Frequency of Statements by Status
plt.figure(figsize=(12,8))

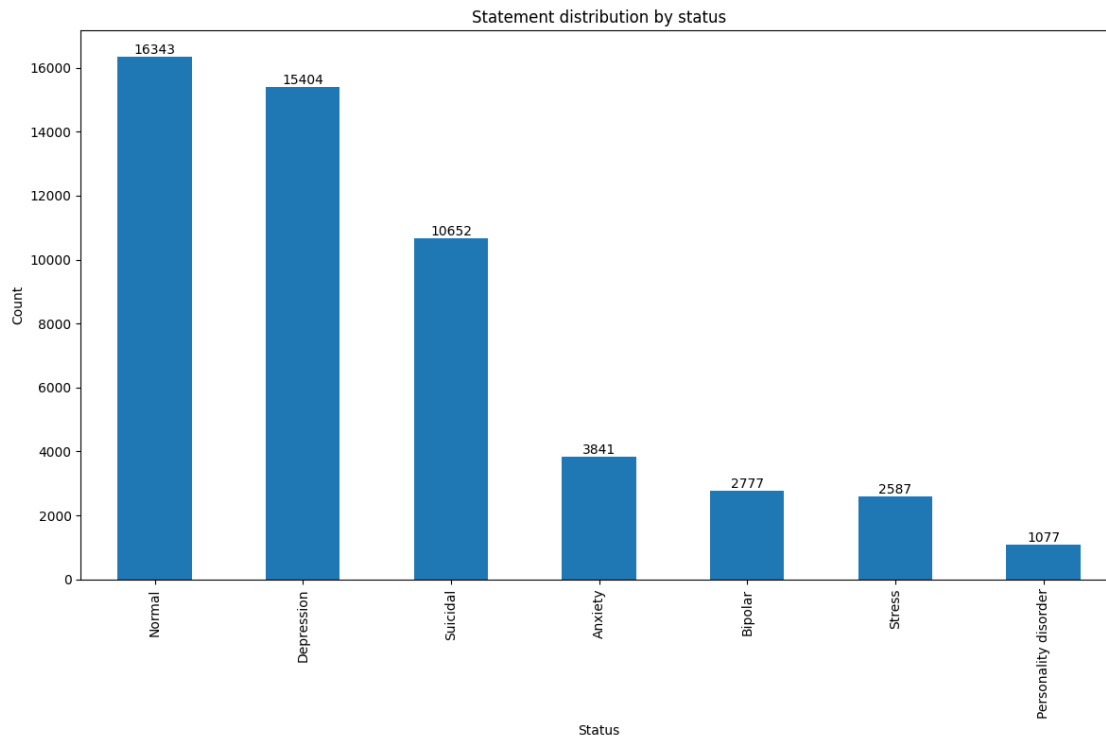
# get the unique status values and their counts
status_counts = df1['status'].value_counts()

# create the bar plot
ax = status_counts.plot(kind='bar')

# add the count labels on top of each bar
for i, v in enumerate(status_counts):
    ax.text(i, v, str(v), ha='center', va='bottom')

plt.title('Statement distribution by status')
plt.xlabel('Status')
plt.ylabel('Count')
```

```
plt.tight_layout()
plt.show()
```



Here is a plot showing distribution by status. Normal is the most common status and contains 16343 data, followed by depression and suicidal, which are the 2nd and 3rd largest portions of the dataset. Personality disorder is the most rare one, which contains 1077 data.

The ratio between different statuses suggests about 70% of the user's input falls under the negative status category.

The target variable in our dataset is unbalanced in favor of depression, suicidal and normal. This imbalance could affect our model's performance, so we'll need to address it later to ensure accurate and fair predictions especially when predicting sentiment analysis for anxiety, bipolar, stress and personality disorder.

```
[17]: # Word Cloud Before Preprocessing
# Combine all statements into a single string
text = ' '.join(df1['statement'].dropna())

# Create a word cloud
wordcloud = WordCloud(width=800, height=400, background_color='white').
    generate(text)

# Display the word cloud
```


10743 ThrowawayIm female, 20 years old. Ever since I...
 10834 I have only 1 person I can somewhat open to bu...
 11537 The title is not meant to discourage others, b...
 11581 I no longer know what else to do but write thi...
 11636 And has life gotten better?​No. Eve...
 11831 Sorry this is long but I doubt anyone will eve...
 13188 I am frustrated. that is the constant theme wi...
 13293 I cannot TAKE IT ANYMORE. I cannot TAKE IT ANY...
 13577 I am very sick and tired, both mentally and ph...
 14602 I am 27 years old and have grown deeper into a...
 16061 Bear with me please, this may be extremely len...
 16498 Hey, this is goodbye note. it is most likely g...
 18215 I am someone living in Turkey. My age is proba...
 18323 I am going to be turning 30 in a couple weeks...
 19321 This happened a little while ago but it still ...
 19701 If there is a more beneficial sub please lmk s...
 20867 Apologies for length. there is a *lot* to expl...
 21285 First I am going to present you with a few que...
 21396 will i ever be noticed? is my life worth anyth...
 21858 I constantly repeat to myself that I have neve...
 22243 I do not expect anyone to read this rambly mes...
 22351 This is a lengthy post but its a summary of my...
 22563 I have been thinking about posting online for ...
 23195 My entire life has spontaneously combusted ove...
 23366 I wish I knew what was wrong with me. So many ...
 23820 I need support or encouragement. I (29M) reall...
 23845 This is a a vent. I (29M) really do not know w...
 24276 I guess it all started when I was I guess 11, ...
 38083 this is my first reddit post also my first tim...
 38255 i m at a very weird place in my life right now...
 38579 hello thank you for reading my post and any ad...
 39579 we ve been seeing a worrying increase in pro s...
 39582 for starter i never really had a childhood whe...
 39752 it doesn t matter anymore i m going to copy an...
 40028 this is a long story i m sorry me and my ex br...
 40208 i m at a very weird place in my life right now...
 40293 i have come to the conclusion that i am just n...
 40371 hello thank you for reading my post and any ad...
 46660 DEPRESSION HAS A PURPOSE: HOW TO USE IT RIGHT ...
 47949 Don't know what to do anymore Back when I was ...
 48915 I think I'm in the middle of a nervous breakdo...
 50253 Manic for 6 months ending up in jail where I h...
 51396 Please help me understand what I went through ...
 52775 I don't know what to do. I don't know how to d...

	status	statement_len
7851	Depression	2153

8221	Depression	1602
9504	Depression	2139
10743	Depression	1537
10834	Suicidal	5248
11537	Depression	2391
11581	Depression	2612
11636	Depression	2415
11831	Depression	2187
13188	Depression	1832
13293	Suicidal	6300
13577	Suicidal	1811
14602	Depression	1809
16061	Depression	1558
16498	Suicidal	1566
18215	Suicidal	2066
18323	Suicidal	1559
19321	Depression	1902
19701	Depression	1661
20867	Depression	1625
21285	Depression	1559
21396	Depression	2510
21858	Depression	2599
22243	Suicidal	2364
22351	Depression	1551
22563	Suicidal	2319
23195	Depression	1818
23366	Depression	1654
23820	Depression	2105
23845	Suicidal	2108
24276	Suicidal	1539
38083	Depression	1559
38255	Depression	1584
38579	Depression	1537
39579	Depression	1747
39582	Depression	1653
39752	Depression	4239
40028	Depression	1726
40208	Depression	1584
40293	Depression	1656
40371	Depression	1537
46660	Bipolar	4727
47949	Depression	1663
48915	Stress	1601
50253	Bipolar	1664
51396	Personality disorder	5419
52775	Anxiety	1586

Many of the longest messages are those with depression and suicidal tendencies. This will help us

since if we shorten the output length when preprocessing the data, we are not reducing the number of data points for those that do not have very many data points such as anxiety, bipolar, stress and personality disorder.

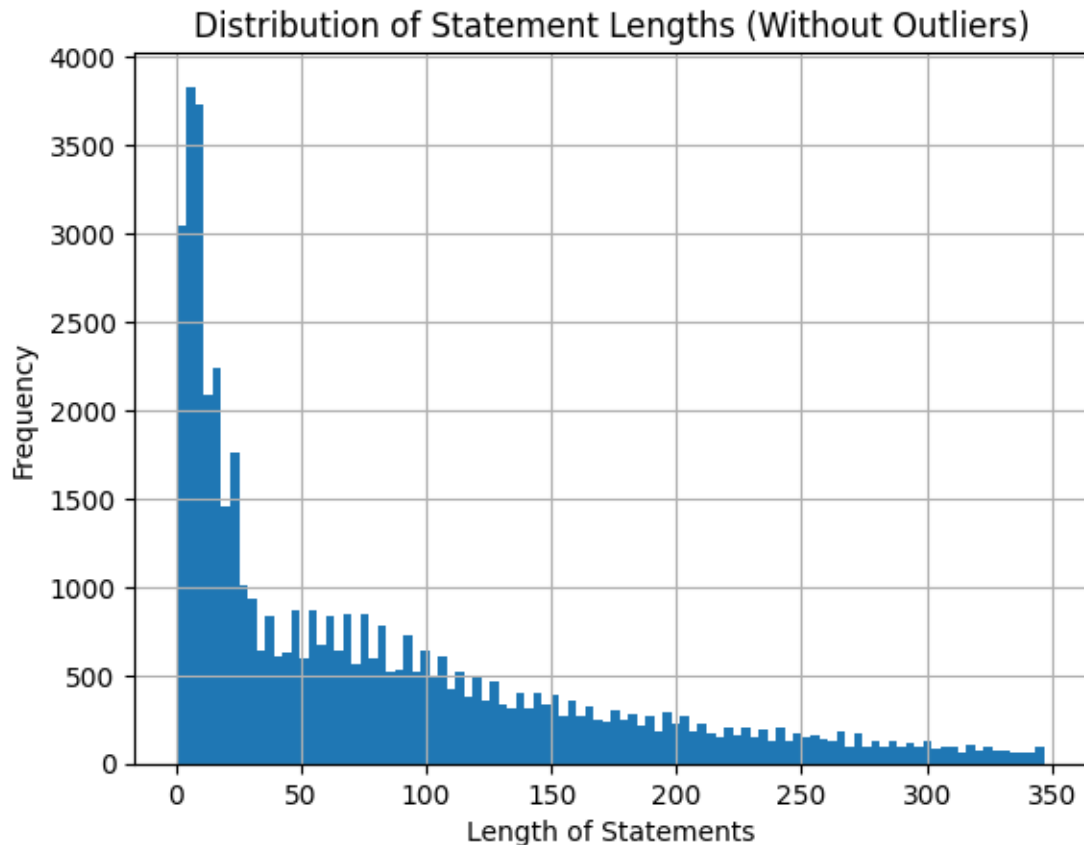
Now we want to see a clearer distribution without these outliers so that we can determine the best output length for preprocessing the text.

```
[19]: # Statement Length Distribution Without Outliers
# Calculate Q1 (25th percentile) and Q3 (75th percentile)
Q1 = df1['statement_len'].quantile(0.25)
Q3 = df1['statement_len'].quantile(0.75)
IQR = Q3 - Q1

# Define the lower and upper bound for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the outliers
filtered_df = df1[(df1['statement_len'] >= lower_bound) & (df1['statement_len'] <= upper_bound)]

# Plot the distribution of statement lengths without outliers
filtered_df['statement_len'].hist(bins=100)
plt.title('Distribution of Statement Lengths (Without Outliers)')
plt.xlabel('Length of Statements')
plt.ylabel('Frequency')
plt.show()
```



This distribution still shows a right-skewed data distribution. We now have a much clearer distribution where approximately 50% of the statements have 0-50 word lengths, especially with a spike at approximately 25 words with approximately 3700 statements. This will help us immensely to determine the best statement length to run our transformer models to save computational resources and time but not decrease model performance.

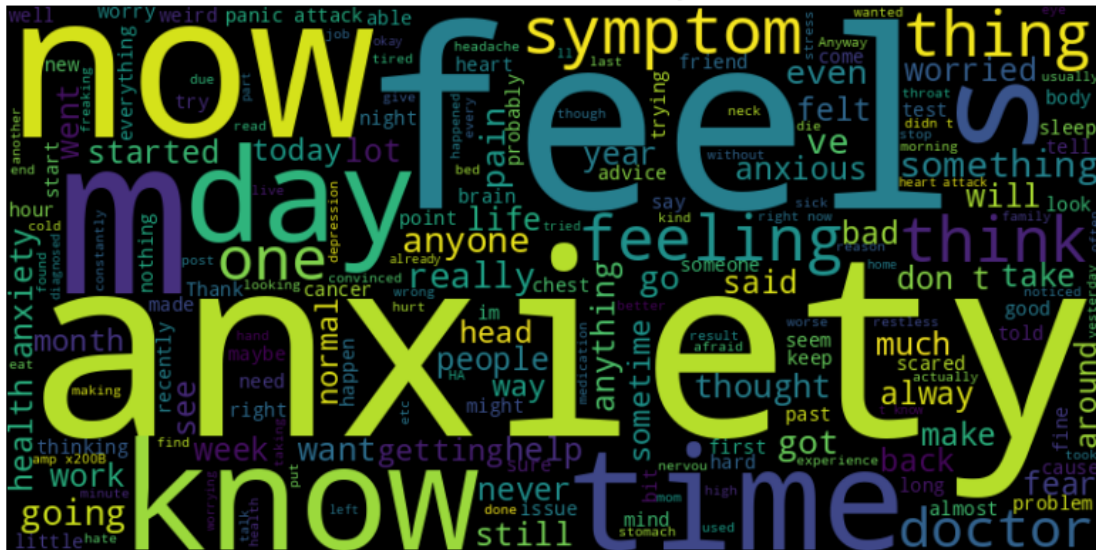
Now we want to take a closer look at the word clouds for each status since it will give us even more information about the possible word indicators for each status.

```
[20]: # Create a function to generate and display a word cloud
def generate_word_cloud(text, title):
    wordcloud = WordCloud(width=800, height=400).generate(text)
    plt.figure(figsize=(10, 5))
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.title(title)
    plt.axis('off')
    plt.show()

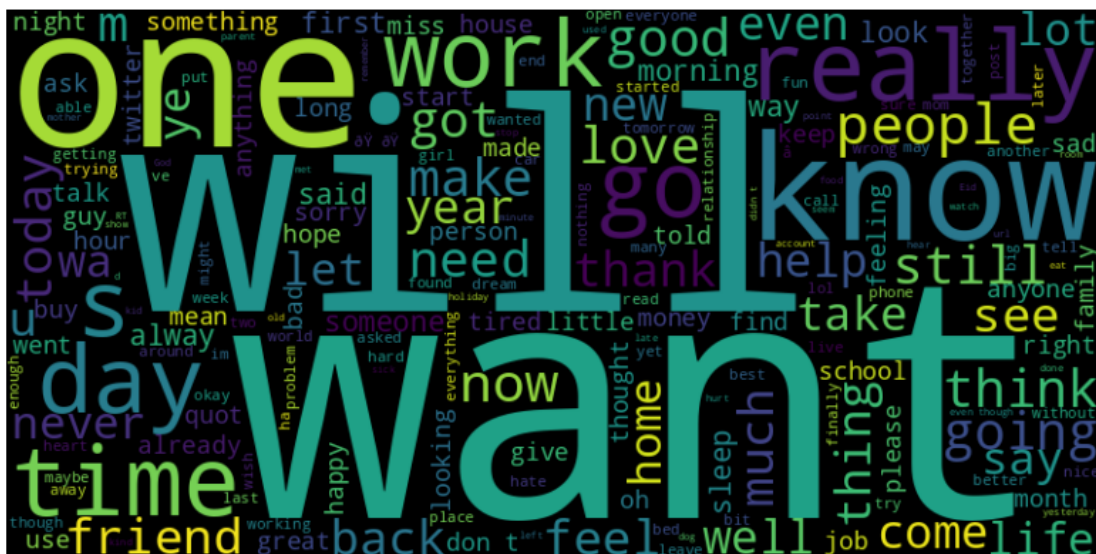
# Generate word clouds for each status
statuses = df1['status'].unique()
```

```
for status in statuses:
    status_text = ' '.join(df1[df1['status'] == status]['statement'])
    generate_word_cloud(status_text, title=f'Word Cloud for {status}')
```

Word Cloud for Anxiety



Word Cloud for Normal



[illegible]

[illegible]

[illegible]

“Will”, “want”, “know” are the most common words for normal. “Life”, “feel”, “want” are the most common words for suicidal. Many people with suicidal tendencies tend to talk about their lives. The words for “normal” status tend to have positive connotations especially “want” and “will”, especially for a “will” to live.

We want to conduct bi-grams and tri-grams analysis for these reasons: Contextual Insights: Bi-grams and tri-grams capture phrases and context that single words (unigrams) might miss. This is particularly important in mental health, where phrases like “feeling down” or “very anxious” provide more insight than individual words.

Identifying Common Themes: Visualizing bi-grams and tri-grams helps identify common themes and expressions in the dataset. This can reveal patterns in how people express their mental health experiences.

16


```

# Tokenization and N-gram generation
# Create a CountVectorizer object with ngram_range set to (2, 3) to generate
↳ bi-grams and tri-grams
vectorizer = CountVectorizer(ngram_range=(2, 3))

# Fit and transform the 'statement' column of the DataFrame to generate the
↳ n-grams
X = vectorizer.fit_transform(df1['statement'])

# Frequency distribution
# Sum the occurrences of each n-gram across all documents
sum_words = X.sum(axis=0)

# Create a list of tuples where each tuple contains an n-gram and its
↳ corresponding frequency
words_freq = [(word, sum_words[0, idx]) for word, idx in vectorizer.vocabulary_.
↳ items()]

# Sort the list of tuples by frequency in descending order
words_freq = sorted(words_freq, key=lambda x: x[1], reverse=True)

# DataFrame for visualization
# Convert the list of tuples into a DataFrame for easier visualization
df_freq = pd.DataFrame(words_freq, columns=['N-gram', 'Frequency'])

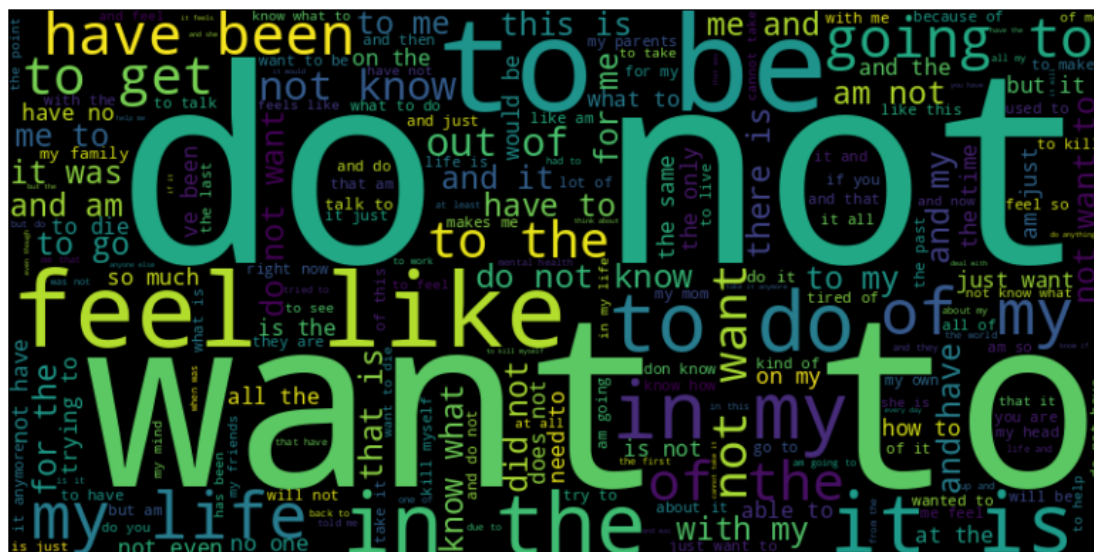
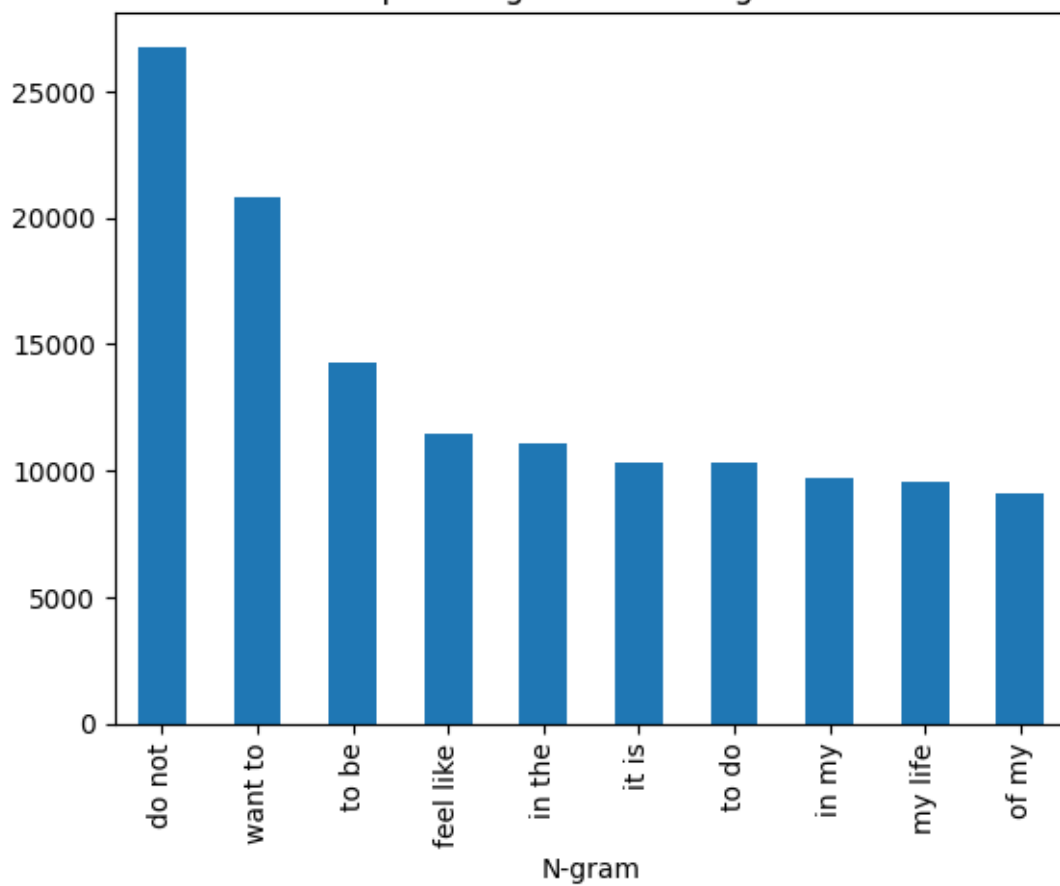
# Bar plot
# Plot the top 10 most frequent n-grams as a bar plot
df_freq.head(10).plot(kind='bar', x='N-gram', y='Frequency', legend=False)
plt.title('Top 10 Bi-grams and Tri-grams')
plt.show()

# Word cloud
# Generate a word cloud from the n-gram frequencies
wordcloud = WordCloud(width=800, height=400).
↳ generate_from_frequencies(dict(words_freq))

# Display the word cloud
plt.figure(figsize=(10, 5))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()

```

Top 10 Bi-grams and Tri-grams



This visualizes the most common bi-grams and tri-grams in our dataset, providing insights into common phrases and patterns, which is particularly useful for sentiment analysis in mental health. It helps identify key expressions and themes that might indicate different emotional states or communication gaps. The top ten are do not, want to, to be, feel like, in the, it is, to do, in my, my life, and of my.

Negative Sentiments: Phrases like “do not” and “feel like” might indicate negative sentiments or expressions of reluctance and emotional states. These bi-grams can help identify statements where individuals are expressing dissatisfaction or discomfort.

Desires and Intentions: Bi-grams such as “want to” and “to do” suggest expressions of desires, intentions, or plans. Analyzing these can reveal what individuals are striving for or what actions they are considering, which can be linked to their mental state.

Self-Reflection: Phrases like “in my,” “my life,” and “of my” indicate self-reflection and personal experiences. These bi-grams can help identify statements where individuals are discussing their personal lives and feelings, which are critical for understanding their mental health.

General Context: Bi-grams like “to be,” “in the,” and “it is” provide general context and can be part of various expressions. While they might not directly indicate sentiment, they help in understanding the structure and flow of the text.

2.0.2 Week 4 Code- Preprocessing the data

df1 is the dataframe that does not have any of the missing values. filtered_df is the dataframe with the outliers removed and no missing values. We will keep the outliers since sentiment analysis often has extreme reviews (e.g., very short or long ones) that can hold strong emotions, valuable for classification especially when using non-transformer models. For transformer models, we may remove outliers by shortening the word length input or use filtered_df to reduce computational time without sacrificing performance. For now, we will use df1 to preprocess the data. Transformer models tend to have different preprocessing techniques anyways.

Warning: After installing imbalanced-learn, please **restart the kernel** for the changes to take effect.

You can do this in Jupyter Notebook by clicking:

Kernel → Restart Kernel

```
[22]: # install packages
      !pip install imbalanced-learn
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: imbalanced-learn in /home/jupyter-
geean/.local/lib/python3.12/site-packages (0.13.0)
Requirement already satisfied: numpy<3,>=1.24.3 in
/opt/tljh/user/lib/python3.12/site-packages (from imbalanced-learn) (1.26.4)
Requirement already satisfied: scipy<2,>=1.10.1 in
/opt/tljh/user/lib/python3.12/site-packages (from imbalanced-learn) (1.13.1)
Requirement already satisfied: scikit-learn<2,>=1.3.2 in
/opt/tljh/user/lib/python3.12/site-packages (from imbalanced-learn) (1.6.1)
Requirement already satisfied: sklearn-compat<1,>=0.1 in /home/jupyter-
```

```

geean/.local/lib/python3.12/site-packages (from imbalanced-learn) (0.1.3)
Requirement already satisfied: joblib<2,>=1.1.1 in
/opt/tljh/user/lib/python3.12/site-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in
/opt/tljh/user/lib/python3.12/site-packages (from imbalanced-learn) (3.5.0)

```

```

[38]: # import packages
import seaborn as sns

import re
import random
from imblearn.over_sampling import RandomOverSampler # used for oversampling
from scipy.sparse import hstack, csr_matrix # To combine sparse matrices

import nltk # nlp package
from nltk.tokenize import word_tokenize # tokenize the text
from nltk.stem import PorterStemmer # stem the text

from sklearn.feature_extraction.text import TfidfVectorizer # used for tf-idf
from sklearn.model_selection import GridSearchCV # used to find the best
    ↪ parameters
from sklearn.preprocessing import LabelEncoder # used to encode categorical
    ↪ variables

```

```

[24]: # Remove duplicates based on 'statement'
df_unique = df1.drop_duplicates(subset=['statement'])

df_unique.head()

```

```

[24]:

```

	statement	status	statement_len
0	oh my gosh	Anxiety	3
1	trouble sleeping, confused mind, restless hear...	Anxiety	10
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14
3	I've shifted my focus to something else but I'...	Anxiety	11
4	I'm restless and restless, it's been a month n...	Anxiety	14

2.1 Why Add Sentence & Character Length in NLP?

Enhancing NLP models like **Naïve Bayes** and **XGBoost** with numerical metadata (e.g., sentence length, character count) improves performance by capturing structural insights.

2.1.1 Benefits:

- **Structural Insights**
 - **Sentence Length:** Differentiates concise vs. verbose texts (e.g., tweets vs. articles).
 - **Character Length:** Indicates complexity, verbosity, or spam tendencies.
- **Better Model Interpretability**

- Helps tree-based models (e.g., XGBoost) make effective splits.
- Useful for readability assessment, spam detection, and authorship identification.
- **Performance Boost**
 - Combines well with TF-IDF, embeddings, and n-grams.
 - Provides independent signals, improving classification accuracy.

2.1.2 Key Use Cases:

- **Sentiment Analysis:** Short reviews are often more direct (positive/negative).

```
[25]: # Calculate the number of characters and sentences

df_unique['num_of_characters'] = df_unique['statement'].str.len()
df_unique['num_of_sentences'] = df_unique['statement'].apply(lambda x: len(nltk.
    ↪sent_tokenize(x)))

# Generate descriptive statistics
description = df_unique[['num_of_characters', 'num_of_sentences']].describe()

# Display the descriptive statistics
print(description)
```

	num_of_characters	num_of_sentences
count	51073.000000	51073.000000
mean	575.375051	6.249251
std	847.661079	10.762749
min	2.000000	1.000000
25%	79.000000	1.000000
50%	313.000000	3.000000
75%	745.000000	8.000000
max	32759.000000	1260.000000

2.2 Text Preprocessing

2.2.1 Convert to lowercase for uniformity

```
[27]: # convert to lowercase
#rename columns
df_unique.rename(columns={'statement': 'original_statement'}, inplace=True)

# create a new cleaned statement column called statement
df_unique['statement']=df_unique['original_statement'].str.lower()
# see the first ten rows
df_unique.head()
```

```
[27]:
```

	original_statement	status	statement_len	\
0	oh my gosh	Anxiety	3	

1	trouble sleeping, confused mind, restless hear...	Anxiety	10
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14
3	I've shifted my focus to something else but I'...	Anxiety	11
4	I'm restless and restless, it's been a month n...	Anxiety	14

	num_of_characters	num_of_sentences	\
0	10	1	
1	64	2	
2	78	2	
3	61	1	
4	72	2	

	statement
0	oh my gosh
1	trouble sleeping, confused mind, restless hear...
2	all wrong, back off dear, forward doubt. stay ...
3	i've shifted my focus to something else but i'...
4	i'm restless and restless, it's been a month n...

2.2.2 Remove URLs and other text, punctuation, and special text

To remove specific patterns such as URLs or other unwanted text (like [View Poll](https://www.reddit.com/poll/...)) from a column in a pandas DataFrame, we can use regular expressions with the re module or pandas built-in string methods. This will help us get cleaner relevant text.

```
[28]: def remove_patterns(text):
    # Remove URLs
    text = re.sub(r'http[s]?://\S+', '', text)
    # Remove markdown-style links
    text = re.sub(r'\[.*?\]\(.*?\)', '', text)
    # Remove handles (that start with '@')
    text = re.sub(r'@\w+', '', text)
    # Remove punctuation and other special characters
    text = re.sub(r'[\W\s]', '', text)
    return text.strip()

# Apply the function to the 'statement' column
df_unique['statement'] = df_unique['statement'].apply(remove_patterns)
# see the first ten rows
df_unique.head()
```

	original_statement	status	statement_len	\
0	oh my gosh	Anxiety	3	
1	trouble sleeping, confused mind, restless hear...	Anxiety	10	
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14	
3	I've shifted my focus to something else but I'...	Anxiety	11	
4	I'm restless and restless, it's been a month n...	Anxiety	14	

	num_of_characters	num_of_sentences	\
0	10	1	
1	64	2	
2	78	2	
3	61	1	
4	72	2	

	statement
0	oh my gosh
1	trouble sleeping confused mind restless heart ...
2	all wrong back off dear forward doubt stay in ...
3	ive shifted my focus to something else but im ...
4	im restless and restless its been a month now ...

2.2.3 Tokenization

This is when the statements are split into words/tokens. Tokenization is **essential in NLP** as it breaks text into smaller units (tokens), making it **processable by models** like Naïve Bayes, XGBoost, and Transformers.

2.2.4 Key Benefits

- **Structures Raw Text** → Converts unstructured text into a usable format.
- **Boosts Model Performance** → Enables better text classification, sentiment analysis, etc.
- **Handles Language Variability**
 - **Word-based:** "I love NLP" → ["I", "love", "NLP"]
 - **Subword-based (BPE, WordPiece):** "unhappiness" → ["un", "happiness"]
 - **Character-based:** Useful for languages without spaces (e.g., Chinese).
- **Improves Efficiency** → Reduces complexity for machine learning models.

```
[30]: # Apply word_tokenize to each element in the 'statement' column
df_unique['tokens'] = df_unique['statement'].apply(word_tokenize)
# see the first ten rows
df_unique.head()
```

```
[30]:
```

	original_statement	status	statement_len	\
0	oh my gosh	Anxiety	3	
1	trouble sleeping, confused mind, restless hear...	Anxiety	10	
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14	
3	I've shifted my focus to something else but I'...	Anxiety	11	
4	I'm restless and restless, it's been a month n...	Anxiety	14	

	num_of_characters	num_of_sentences	\
0	10	1	
1	64	2	
2	78	2	
3	61	1	
4	72	2	

	statement	\
0	oh my gosh	
1	trouble sleeping confused mind restless heart ...	
2	all wrong back off dear forward doubt stay in ...	
3	ive shifted my focus to something else but im ...	
4	im restless and restless its been a month now ...	

	tokens
0	[oh, my, gosh]
1	[trouble, sleeping, confused, mind, restless, ...]
2	[all, wrong, back, off, dear, forward, doubt, ...]
3	[ive, shifted, my, focus, to, something, else,...]
4	[im, restless, and, restless, its, been, a, mo...]

2.3 What is Stemming & Why is it Important in NLP?

Stemming reduces words to their root form by removing prefixes/suffixes, helping **normalize** text and **reduce dimensionality** in NLP.

2.3.1 How It Works

- "running" → "run"
- "happily" → "happi"
- "flies" → "fli"

2.3.2 Why It Matters

- **Reduces Vocabulary Size** → Groups similar words.
- **Boosts Search & NLP Models** → "run" and "running" treated alike.
- **Speeds Up Processing** → Fewer unique tokens.

2.3.3 Limitation

- Can produce **incorrect roots** ("better" → "bet").
- **Lemmatization** (more accurate) uses a dictionary.

2.3.4 Takeaway

Stemming **simplifies text, reduces redundancy, and improves efficiency** in NLP.

```
[31]: # Initialize the stemmer
stemmer = PorterStemmer()

# Function to stem tokens and convert them to strings
def stem_tokens(tokens):
    return ' '.join(stemmer.stem(str(token)) for token in tokens)

# Apply the function to the 'tokens' column
df_unique['tokens_stemmed'] = df_unique['tokens'].apply(stem_tokens)

# print the first ten rows
df_unique.head()
```

```
[31]:
```

	original_statement	status	statement_len	\
0	oh my gosh	Anxiety		3
1	trouble sleeping, confused mind, restless hear...	Anxiety	10	
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	14	
3	I've shifted my focus to something else but I'...	Anxiety	11	
4	I'm restless and restless, it's been a month n...	Anxiety	14	

	num_of_characters	num_of_sentences	\
0	10	1	
1	64	2	
2	78	2	
3	61	1	
4	72	2	

	statement	\
0	oh my gosh	
1	trouble sleeping confused mind restless heart ...	
2	all wrong back off dear forward doubt stay in ...	
3	ive shifted my focus to something else but im ...	
4	im restless and restless its been a month now ...	

	tokens	\
0	[oh, my, gosh]	
1	[trouble, sleeping, confused, mind, restless, ...	
2	[all, wrong, back, off, dear, forward, doubt, ...	
3	[ive, shifted, my, focus, to, something, else,...	
4	[im, restless, and, restless, its, been, a, mo...	

	tokens_stemmed
0	oh my gosh
1	troubl sleep confus mind restless heart all ou...

```
2 all wrong back off dear forward doubt stay in ...
3 ive shift my focu to someth els but im still w...
4 im restless and restless it been a month now b...
```

2.4 Not Removing Stop Words

Stop words are **frequent words** that may not carry significant meaning in NLP tasks.

2.4.1 General Stop Words

- **Articles** → *a, an, the*
- **Prepositions** → *in, on, at, by, with*
- **Pronouns** → *I, you, he, she, it, they*
- **Conjunctions** → *and, but, or, so*
- **Auxiliary Verbs** → *is, are, was, were, have, do, does*

For **mental health sentiment analysis**, it's best to **keep stop words** because:

2.4.2 Context Matters

- Words like “**not**,” “**never**,” “**very**” can flip sentiment.
 - “*not okay*” “*okay*”

2.4.3 Emotional Expressions

- Stop words are essential for capturing **feelings and emotions**.
 - “*I feel so lost*” carries more meaning than “*feel lost*”.

We can see from the word clouds that there are not many stop words that are in bold so we will not remove them.

```
[ ]: # Now we create our dataset for train-validation-test adding the numerical_
    ↪ features
```

```
[32]: X = df_unique[['tokens_stemmed', 'num_of_characters', 'num_of_sentences']]
      y = df_unique['status']
```

2.5 Why Use Label Encoding for Categorical Variables?

Label encoding converts **categorical target variables (Y)** into numerical format for machine learning models.

2.5.1 Why is it Necessary?

- **ML Models Require Numeric Input** → Algorithms like **XGBoost**, **Naïve Bayes**, **SVM** can't process text labels.

- **Standardizes Target Variable** → Maps categories to integers (e.g., “positive” → 2, “neutral” → 1, “negative” → 0).
- **Compatible with Many Models** → Needed for both **classification** and **regression tasks**.

```
[33]: # label encode our categorical variables for y
      lbl_enc = LabelEncoder()
      y = lbl_enc.fit_transform(y.values)
```

2.5.2 Train-Validation-Test Split

Most common splits are 80-20 so we will use this split here. We will also create a validation set that is 10% and the test set is 10%. The final splits will be 80-10-10.

I did these splits, since this is a common split in machine learning and data science but also because with 80% of the data as a training dataset, a large portion ensures that the model has enough data to learn from, which helps in capturing the underlying patterns and relationships in the data. With a Validation Set of 10%, it can tune hyperparameters and make decisions about the model architecture and helps prevent overfitting by providing a checkpoint to evaluate the model's performance on unseen data during the training process. With a 10% Test Set, we can evaluate the model's performance after it has been trained and validated and 10% is a large enough size given that we have approximately 50,000 data points. The 80-10-10 split is a balanced approach that ensures the model has sufficient data for training while also providing enough data for validation and testing to ensure robust performance

```
[34]: from sklearn.model_selection import train_test_split

      # Split the data into training (80%) and temporary (20%) sets
      train_x, temp_x, train_y, temp_y = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

      # Split the temporary set into validation (50%) and test (50%) sets
      val_x, test_x, val_y, test_y = train_test_split(temp_x, temp_y, test_size=0.5,
      ↪random_state=42)

      # Print the sizes of the splits
      print(f"Training set size: {len(train_x)}")
      print(f"Validation set size: {len(val_x)}")
      print(f"Test set size: {len(test_x)}")
```

```
Training set size: 40858
Validation set size: 5107
Test set size: 5108
```

2.5.3 Convert Text to Features using tf-idf

Now, we will transform tokens (words) into numerical values that represent the importance of words in a document relative to a collection of documents. This helps highlight unique words in a document while downplaying common ones, making it easier for machine learning models to

identify relevant patterns and make better predictions.

```
[39]: # 1. Initialize TF-IDF Vectorizer and fit/transform on the 'tokens_stemmed' column
vectorizer = TfidfVectorizer(ngram_range=(1, 2), max_features=50000)
X_train_tfidf = vectorizer.fit_transform(train_x['tokens_stemmed'])
X_val_tfidf = vectorizer.transform(val_x['tokens_stemmed'])

# 2. Extract numerical features and convert to sparse matrix
X_train_num = csr_matrix(train_x[['num_of_characters', 'num_of_sentences']].values)
X_val_num = csr_matrix(val_x[['num_of_characters', 'num_of_sentences']].values)

# 3. Combine TF-IDF features with numerical features
X_train_combined = hstack([X_train_tfidf, X_train_num])
X_val_combined = hstack([X_val_tfidf, X_val_num])

# Print number of feature words
print('Number of feature words:', len(vectorizer.get_feature_names_out()))
```

Number of feature words: 50000

```
[40]: X_train_combined.shape
```

```
[40]: (40858, 50002)
```

2.6 Why Oversampling for an Imbalanced Dataset?

In an **imbalanced dataset**, the model may favor the majority class, leading to **biased predictions**. **Oversampling** helps balance the dataset by increasing minority class samples.

2.6.1 Why is Oversampling Needed?

- **Prevents Majority Class Bias** → Ensures the model learns patterns from both classes.
- **Improves Model Performance** → Leads to better recall, F1-score, and generalization.
- **Enhances Minority Class Representation** → Avoids underestimating rare but important cases.

2.6.2 Why Does Random Over-Sampling Work Best?

Maintains Original Data Distribution → Simply duplicates minority class samples, avoiding synthetic noise (SMOTE).

Preserves Minority Class Variability → Unlike SMOTE, which may create unrealistic synthetic samples.

Avoids Data Loss → Unlike Under-Sampling, which removes majority class samples and risks losing valuable information.

```
[41]: # Apply Random Over-Sampling on the vectorized data  
      ros = RandomOverSampler(random_state=101)  
      X_train_resampled, y_train_resampled = ros.  
          ↪ fit_resample(X_train_combined, train_y)
```

```
[42]: # see the new dataset  
      X_train_resampled.shape
```

```
[42]: (89215, 50002)
```